

CSE 200 - Technical Writing and Presentation

BANGLADESH UNIVERSITY OF
ENGINEERING AND TECHNOLOGY

DEPARTMENT OF COMPUTER SCIENCE AND
ENGINEERING

DEEPPFAKE DETECTION

Seeing Is No Longer Believing

September 25, 2026

Md. Mufeed Talukder : 151
Tanjimul Islam : 159
Md Tahmidul Islam : 180

Contents

1	Introduction	2
2	Generation of Deepfake Media	3
2.1	General Architecture	3
2.2	Mathematical Formulation	3
2.3	Illustrative Example	4
3	Passive Detection: Finding the Flaws	4
3.1	Spatial Artifact Detection	5
3.2	Temporal Artifact Detection	5
3.3	Multimodal Detection	6
3.4	The Cat-and-Mouse Problem	6
4	Proactive Defense: Proving Provenance	6
4.1	Invisible Watermarking	6
4.2	Content Provenance Standards	7
5	Applications and Defense in Depth	8
5.1	Where Detection Matters	8
5.2	Defense in Depth	8
6	Challenges and Future Directions	9
7	Conclusion	10

Deepfake Detection

Abstract

A deepfake is a synthetic image, video, or audio clip generated by deep neural networks to resemble a real person convincingly. Advances in autoencoder- and Generative Adversarial Network (GAN)-based architectures have made such media cheap to produce and difficult to distinguish from authentic recordings, creating a growing risk to journalism, elections, financial systems, and personal privacy. This report examines the deepfake problem from generation through detection to prevention. We first describe the encoder-decoder pipeline that produces synthetic media and the objective functions that drive it. We then present two complementary detection strategies: passive detection, which inspects spatial, temporal, and audio-visual artifacts left behind by the generation process, and proactive defense, which embeds cryptographic watermarks and provenance metadata at the point of creation so that origin can be verified later. We compare the two strategies, outline where each is used in practice, and argue that no single detector is a durable solution because detection and generation are locked in an adversarial cycle. We close by discussing the shift toward content provenance standards such as C2PA and the open challenges that remain.

1 Introduction

A **deepfake** is a piece of synthetic media, typically an image, a video, or an audio clip, in which a deep neural network has replaced or altered a person’s likeness so that the result appears authentic. The term is a portmanteau of “deep learning” and “fake,” and it reflects the underlying technique: a neural network learns the statistical structure of a person’s face or voice from real examples and then uses that structure to generate new, fabricated content.

Deepfake technology is not inherently malicious. The same generative models are used for film dubbing, historical restoration, and accessibility tools. However, the ease with which convincing fake media can now be produced has created a genuine problem of trust. A fabricated speech by a public official, a cloned voice used to authorize a fraudulent wire transfer, or a manipulated video used to harass an individual are no longer hypothetical scenarios; each has already occurred. Because a single realistic fake can spread faster than a correction, the cost of being wrong about authenticity has grown substantially.

This report is organized around three questions that mirror the stages of the deepfake life cycle.

- How is synthetic media generated? (Section 2)
- How can synthetic media be detected after it has been created? (Section 3)

- How can the origin of media be proven before it is ever questioned? (Section 4)

We conclude by comparing these approaches, listing domains where detection is already deployed, and outlining open challenges for the field (Sections 5–7).

2 Generation of Deepfake Media

Before detection can be discussed meaningfully, we must understand what a detector is looking for. This section describes the general architecture used to generate face- and voice-based deepfakes and the mathematical objective that drives training.

2.1 General Architecture

Most face-swapping and face-reenactment systems share a common structure, illustrated in Figure 1. Two inputs are required: a *source identity*, which supplies the face or voice to be inserted, and a *target context*, which supplies the video or audio in which that identity should appear. Both inputs are passed to a deep learning engine, most commonly an autoencoder or a Generative Adversarial Network (GAN), which encodes each input into a compact latent representation and then decodes that representation into a synthetic output that combines the identity of the source with the motion, lighting, and background of the target.




Figure 1: General deepfake generation pipeline. A source identity and a target context are encoded into a shared latent space and decoded into a synthetic output.

2.2 Mathematical Formulation

Let x denote an input image, video frame, or audio segment drawn from a training dataset $\mathcal{D}$. An autoencoder-based generator consists of an encoder E_ϕ and a decoder G_θ , where ϕ and θ are learned parameters. The encoder maps the input to a latent vector $z = E_\phi(x)$, and the decoder reconstructs an output $\hat{x} = G_\theta(z)$. Training minimizes the reconstruction loss

$$\mathcal{L}_{\text{recon}} = \mathbb{E}_{x \sim \mathcal{D}} [\|x - \hat{x}\|_2^2], \quad (1)$$

where $\|\cdot\|_2$ denotes the Euclidean norm. A face-swapping system typically trains one shared encoder alongside two separate decoders, one for each identity, so that a face encoded from person A can be decoded through person B’s decoder to produce a face that has A’s expression and B’s identity.

Many modern generators improve realism further by pairing the decoder G_θ with a discriminator D_ψ , forming a Generative Adversarial Network. The discriminator is trained

to distinguish real data from generated data, while the generator is trained to fool the discriminator. This adversarial relationship is expressed as the minimax objective

$$\min_{\theta} \max_{\psi} V(D_{\psi}, G_{\theta}) = \mathbb{E}_{x \sim \mathcal{D}} [\log D_{\psi}(x)] + \mathbb{E}_{z \sim p(z)} [\log (1 - D_{\psi}(G_{\theta}(z)))], \quad (2)$$

where $p(z)$ is a prior distribution, often Gaussian, from which latent vectors are sampled. As training proceeds, D_{ψ} becomes better at spotting fakes, which in turn forces G_{θ} to produce increasingly realistic output; this adversarial pressure is precisely what makes GAN-based deepfakes difficult to detect by eye. It is also the same adversarial pattern, generator against discriminator, that reappears later at the level of the entire field, where a detector plays the role of D_{ψ} and a new generation technique plays the role of G_{θ} (Section 3.4).

2.3 Illustrative Example

Figure 2 shows a single face-swap example that follows the pipeline of Figure 1: a source image supplies the identity, a target image supplies the pose and background, and the outcome is a synthetic frame in which the source identity has been transferred onto the target context. The example reused from our presentation slides is suitable for this figure; no new image is required here.



Figure 2: Source, target, and outcome for a single face-swap example (reuse `example_source`, `example_target`, `example_output` from the presentation).

3 Passive Detection: Finding the Flaws

Detection strategies fall into two broad categories, summarized in Table 1: *passive detection*, which examines media after it has been created, and *proactive defense*, which embeds safeguards into the generation process itself. This section covers passive detection; proactive defense is discussed in Section 4.

Aspect	Passive Detection	Proactive Defense
When it acts	After media has been created	At the moment media is created
Method	Examines media for hidden artifacts left by the generation process	Embeds a cryptographic watermark or provenance record into the media
Posture	Reactive	Preventative
Typical question	“Does this look fake?”	“Can we prove where this came from?”

Table 1: Comparison of passive detection and proactive defense.

3.1 Spatial Artifact Detection

A single frame is often enough to reveal a deepfake. Generative models tend to leave two characteristic spatial artifacts. First, skin regions are frequently over-smoothed because the decoder averages over many training examples, producing a waxy texture that lacks natural pore-level detail. Second, generated faces often exhibit an unnaturally even focus: the face and the background are equally sharp, whereas a real photograph taken with a camera lens shows a shallow depth of field that blurs the background relative to the subject.

These cues can be learned automatically. A Convolutional Neural Network (CNN) is trained as a binary classifier $f_w : x \mapsto p \in [0, 1]$, where x is an input frame, w denotes the network’s weights, and p is the predicted probability that the frame is synthetic. Given a labeled training set of real and fake frames with ground-truth labels $y \in \{0, 1\}$, the network is trained by minimizing the binary cross-entropy loss

$$\mathcal{L}_{\text{BCE}} = -\frac{1}{N} \sum_{i=1}^N [y_i \log(p_i) + (1 - y_i) \log(1 - p_i)], \quad (3)$$

where N is the number of training examples in a mini-batch. Figure 3 illustrates the two artifacts described above together with the classification pipeline.




Figure 3: A CNN classifier trained to distinguish real frames from synthetic ones using spatial cues such as skin smoothness and depth-of-field consistency.

3.2 Temporal Artifact Detection

A single frame can look convincing while the video as a whole does not, because generative models operate mostly on a frame-by-frame basis and are not always consistent across time. Three temporal cues are commonly exploited.

- **Blink rate:** humans blink at a fairly regular interval; early face-swap datasets under-represented closed-eye frames, causing generated subjects to blink too rarely or too regularly.
- **Motion smoothness:** synthetic faces can exhibit subtle jitter or warping between consecutive frames that a real camera recording would not produce.

- **Physiological signal:** a real face shows faint, periodic color shifts caused by blood flow beneath the skin (remote photoplethysmography); this signal is difficult for a generator to reproduce consistently.

Because these cues depend on the relationship between frames rather than on any single frame, they are best captured by a sequence model such as a Long Short-Term Memory (LSTM) network or a Vision Transformer applied across a short clip $\{x_1, x_2, \dots, x_T\}$, where T is the number of frames in the clip. The sequence model is trained with the same binary cross-entropy objective as Equation 3, but the input is now a temporal window rather than a single frame.

3.3 Multimodal Detection

Voice can be cloned just as convincingly as a face can be swapped, and an attacker who fakes only one modality often leaves the other unchanged, which creates an inconsistency that can be measured directly. The most reliable cue in this category is lip-sync mismatch: the shape of the mouth in the video should match the phoneme being produced in the audio at every instant, and a mismatch of even a few frames is detectable even when neither the video nor the audio is individually flagged as fake by a single-modality detector.

A multimodal detector processes the video stream and the audio stream separately, using a video encoder for the former and an audio encoder in the style of Whisper or wav2vec2 for the latter, and then combines the two streams with a cross-modal transformer that checks whether the acoustic and visual signals are consistent with each other. This cross-checking approach is powerful precisely because it does not require either modality to be individually recognizable as synthetic; it only requires the two to disagree.

3.4 The Cat-and-Mouse Problem

Passive detection faces a structural limitation. Every new detection technique that becomes public gives generator designers a specific artifact to correct in the next model version, so detection accuracy that is measured on today’s fakes routinely degrades on tomorrow’s. This dynamic mirrors the adversarial training objective of Equation 2 at the scale of the research community: detectors act as an external discriminator, and generators are iteratively improved to defeat them. Because detection is always one step behind the latest generation technique, passive detection alone cannot be a permanent solution; it must be paired with a proactive strategy, which we discuss next.

4 Proactive Defense: Proving Provenance

Rather than trying to spot a fake after it exists, a proactive defense makes authenticity verifiable from the moment media is created.

4.1 Invisible Watermarking

An invisible watermark is a pattern embedded directly into a generated image or video at creation time, using a cryptographic key k known to the issuing organization. The

watermark embedding function can be described as $x_w = W_k(x)$, where x is the generator’s raw output, x_w is the watermarked output, and W_k is the embedding operator. The pattern is designed to be perceptually invisible, so x_w and x look identical to a human viewer, while remaining statistically detectable to a verification system that has access to the key. A well-designed watermark is also robust: it should survive common transformations such as compression, cropping, and re-encoding through a screenshot, so that the provenance signal is not lost as content is shared across platforms.

Verification is the mirror operation of embedding: given a piece of media and the same key k , a verification system extracts a candidate pattern and checks it for a match. Figure 4 summarizes the pipeline from generation through verification.



Figure 4: Invisible watermarking pipeline: a key-dependent pattern is embedded at generation time and later extracted by a verification system.

4.2 Content Provenance Standards

Watermarking answers a narrow question: was this specific pattern present? A broader approach, known as content provenance, attaches a signed record of an asset’s history, its origin, its creator, and every edit it has undergone, directly to the file as metadata. The Coalition for Content Provenance and Authenticity (C2PA) is the leading industry standard in this space; it defines *Content Credentials* that travel with a media file and can be inspected by any compliant viewer. Table 2 contrasts this provenance-based approach with the detection-based approach of Section 3.

Aspect	Detection	Provenance
Question asked	Does this look fake?	Can we prove where this came from?
Basis	Statistical inference over artifacts	Cryptographic signature and metadata
Confidence	Probabilistic; can be wrong	Deterministic, if the signature is intact
Weakness	Loses ground as generators improve	Requires adoption across creation tools and platforms

Table 2: Detection asks whether content looks fake; provenance asks whether its origin can be proven.

5 Applications and Defense in Depth

5.1 Where Detection Matters

The guarantees required of a detection system differ sharply by domain, as summarized in Table 3.

Domain	Typical Use	Primary Requirement
Social platforms	Upload scanning, content labeling	High throughput, low latency
FinTech and identity	Fraud prevention, liveness checks	Low false-negative rate
Courts and forensics	Evidence auditing, forensic verification	High explainability, legal defensibility
Government and elections	Statement verification, media integrity	Public trust, rapid response

Table 3: Representative application domains for deepfake detection and their primary requirements.

5.2 Defense in Depth

No single technique from Sections 3 and 4 is sufficient on its own. In practice, an effective system layers three stages, illustrated in Figure 5: *prevention*, which limits access to generation tools and embeds watermarks; *detection*, which screens content already in circulation using spatial, temporal, and audio-visual cues; and *verification*, which checks cryptographic provenance when it is available. Each stage compensates for the weaknesses of the others, so that a failure at one stage does not compromise the whole system.




Figure 5: Defense in depth: prevention, detection, and verification act as three complementary layers rather than a single point of failure.

Algorithm 1 summarizes the operational logic of a layered detection service that combines a passive classifier with provenance verification.

Algorithm 1 Layered Screening of an Uploaded Media Item

Require: Media item x , watermark key k , spatial detector f_w , temporal detector g_v , decision threshold τ

- 1: Attempt to extract a provenance record or watermark from x using key k
- 2: **if** a valid, matching record is found **then**
- 3: **return** VERIFIED AUTHENTIC
- 4: **else**
- 5: Compute spatial score $p_s \leftarrow f_w(x)$ (Equation 3)
- 6: Compute temporal score $p_t \leftarrow g_v(x)$ over a short clip
- 7: Combine scores: $p \leftarrow \frac{1}{2}(p_s + p_t)$
- 8: **if** $p > \tau$ **then**
- 9: **return** FLAGGED AS LIKELY SYNTHETIC
- 10: **else**
- 11: **return** PASSED, NO PROVENANCE RECORD
- 12: **end if**
- 13: **end if**

6 Challenges and Future Directions

Despite the progress surveyed above, several open challenges limit how far current detection and provenance methods can be trusted in practice. Each is paired below with the direction the field is taking in response.

- **Poor generalization across generators.** Detectors trained on one generator family generalize poorly to another, because each generation method leaves a different artifact signature.
Response: Multimodal detectors that fuse image, video, and audio cues, rather than relying on any single artifact.
- **Accuracy decay over time.** Detection accuracy measured today decays as new generation techniques are released (Section 3.4).
Response: Continuous, ongoing verification rather than a one-time trained model.
- **Incomplete provenance adoption.** Watermarking and provenance require every creation tool and platform to participate before they are reliable.
Response: Standardized provenance by default, led by initiatives such as C2PA.
- **Risk of full automation.** Fully automated decisions carry legal and ethical risk in high-stakes domains such as courts and elections.
Response: Human-in-the-loop systems, where the model flags content and a person makes the final judgment.

Taken together, these directions point toward detection systems that are multimodal, continuously updated, standardized at the point of creation, and paired with human judgment in consequential decisions — rather than a single static classifier deployed once and trusted indefinitely.

7 Conclusion

Deepfake generation and deepfake detection are two sides of the same adversarial process described by Equation 2: as generators improve, detectors must improve to match them, and the cycle repeats. We have shown that passive detection, whether based on spatial artifacts (Section 3.1), temporal inconsistency (Section 3.2), or cross-modal mismatch (Section 3.3), remains useful but is structurally unable to win the race on its own, because it always responds to a generation technique that already exists. Proactive defense through watermarking and content provenance (Section 4) changes the question from “does this look fake?” to “can we prove where this came from?”, and that shift is what makes verification durable even as generation techniques continue to advance.

The goal of this line of research is not a perfect detector; no fixed classifier can remain accurate indefinitely against a generator trained specifically to defeat it. The realistic goal is a layered ecosystem, combining prevention, detection, and verification (Section 5.2), in which fabricated media is harder to spread convincingly and genuine media is easier to verify. As the underlying technology matures, we expect the balance to continue shifting from after-the-fact detective work toward built-in, cryptographically provable authenticity.

References

- [1] I. J. Goodfellow, J. Pouget-Abadie, M. Mirza, B. Xu, D. Warde-Farley, S. Ozair, A. Courville, and Y. Bengio, “Generative Adversarial Networks,” *Advances in Neural Information Processing Systems (NeurIPS)*, 2014.
- [2] A. Rössler, D. Cozzolino, L. Verdoliva, C. Riess, J. Thies, and M. Nießner, “FaceForensics++: Learning to Detect Manipulated Facial Images,” *International Conference on Computer Vision (ICCV)*, 2019.
- [3] Y. Li, X. Yang, P. Sun, H. Qi, and S. Lyu, “Celeb-DF: A Large-Scale Challenging Dataset for DeepFake Forensics,” *IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2020.
- [4] S. Agarwal, H. Farid, O. Fried, and M. Agrawala, “Detecting Deep-Fake Videos from Phoneme-Viseme Mismatches,” *CVPR Workshops*, 2020.
- [5] Coalition for Content Provenance and Authenticity (C2PA), *C2PA Technical Specification*, <https://c2pa.org>, accessed 2026.